

Image Classification Using Deep Learning on CIFAR-10 Dataset

Gajan Mohan Raj

TABLE OF CONTENTS

01

Background and
Requirements

02

Project
Approach

03

Findings

04

NAME OF THE
SECTION

01

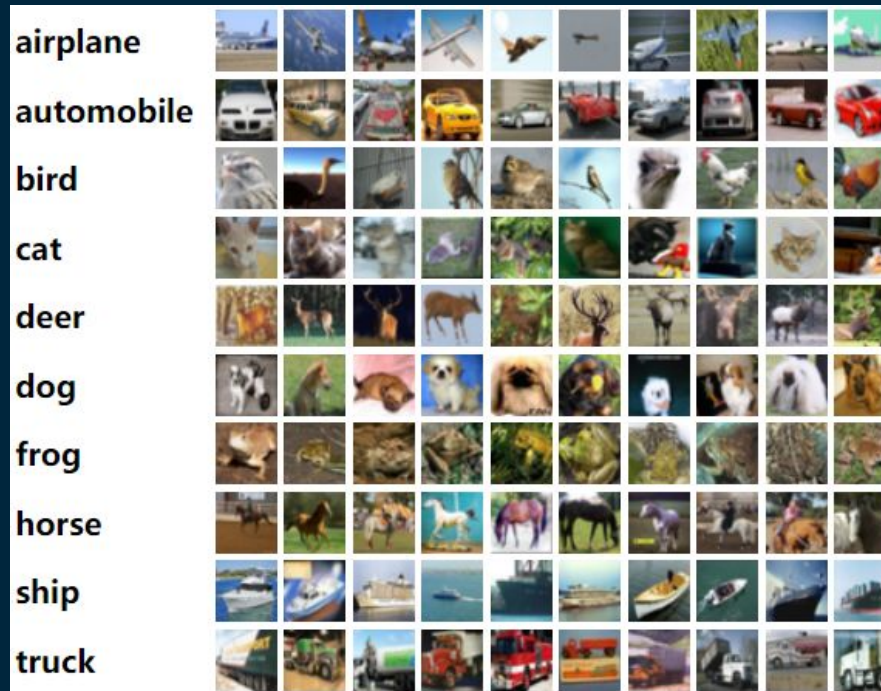
Background and Requirements

Background and Requirements

In this project, I was tasked with building and training a Convolutional Neural Network (CNN) to classify images from the CIFAR-10 dataset. CIFAR-10 is a well-known dataset used in the machine learning community, consisting of 60,000 32x32 color images in 10 different classes.

Requirements:

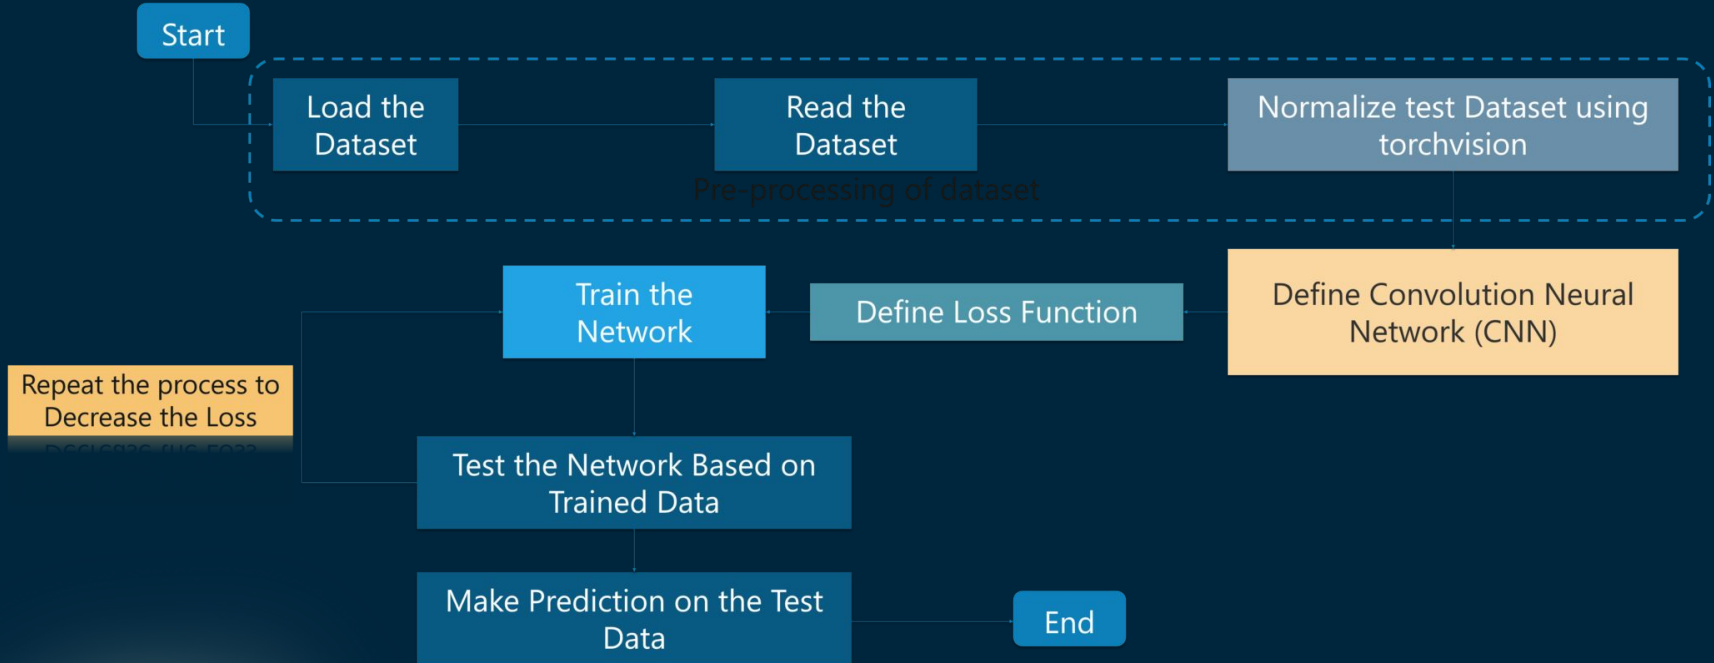
- Programming Language: Python
- Main Library: PyTorch Additional Libraries:
- Torchvision: For handling image loading and transformations.
- Matplotlib: For visualizing images and training results.



02

Project Approach

Project Approach: General Layout



Project Approach: Building the CNN Architecture

Convolutional neural networks (CNNs) are commonly used for image classification.

In this project, I constructed a CNN architecture with a series of convolutional blocks followed by dense (fully-connected) layers:

- Several convolutional blocks, where each block typically includes a convolutional layer followed by batch normalization and ReLU activation
- Max pooling layers placed periodically between convolutional blocks to reduce spatial dimension of feature maps and control overfitting
- Dense layers include batch normalization and dropout layers for regularization

Layer (type)	Output Shape	Param #
Conv2d-1	[-1, 32, 32, 32]	896
BatchNorm2d-2	[-1, 32, 32, 32]	64
Conv2d-3	[-1, 64, 32, 32]	18,496
BatchNorm2d-4	[-1, 64, 32, 32]	128
MaxPool2d-5	[-1, 64, 16, 16]	0
Conv2d-6	[-1, 128, 16, 16]	73,856
BatchNorm2d-7	[-1, 128, 16, 16]	256
Conv2d-8	[-1, 128, 16, 16]	147,584
BatchNorm2d-9	[-1, 128, 16, 16]	256
MaxPool2d-10	[-1, 128, 8, 8]	0
Conv2d-11	[-1, 256, 8, 8]	295,168
BatchNorm2d-12	[-1, 256, 8, 8]	512
Conv2d-13	[-1, 256, 8, 8]	590,080
BatchNorm2d-14	[-1, 256, 8, 8]	512
MaxPool2d-15	[-1, 256, 4, 4]	0
Linear-16	[-1, 1024]	4,195,328
BatchNorm1d-17	[-1, 1024]	2,048
Dropout-18	[-1, 1024]	0
Linear-19	[-1, 512]	524,800
BatchNorm1d-20	[-1, 512]	1,024
Dropout-21	[-1, 512]	0
Linear-22	[-1, 10]	5,130


Project Approach: Training Process

- Dataset preparation

- There are 60,000 images of the CIFAR-10 dataset
- All images were first normalized to improve performance and stability of the CNN
- Then split into three distinct datasets using `random_split`
 - Train set
 - 40,000 images
 - Validation set
 - 10,000 images
 - Test set
 - 10,000 images

- Hyperparameters



- *Loss function*: Cross entropy loss
- *Optimizer*: Adam optimizer
- *Epochs*: 15
- *Learning rate*: 0.001
- *Batch size*: 64



03


Findings

Findings: Training and Validation Curves for Original CNN



- The CNN architecture performed very well with an accuracy of **83.90%** on the validation set
- The validation loss tended to increase after the 7th epoch while the training loss continued to decrease, indicating overfitting
 - Early stopping and extra regularization techniques could be used to prevent overfitting

Findings: Performance on Test Set



- The CNN architecture performed very well on the test set with an impressive accuracy of **84%**!

Classification Report				
	Precision	Recall	F1-Score	Support
Airplane	0.83	0.88	0.85	1000
Car	0.88	0.95	0.91	1000
Bird	0.80	0.73	0.76	1000
Cat	0.70	0.70	0.70	1000
Deer	0.84	0.81	0.83	1000
Dog	0.77	0.79	0.78	1000
Frog	0.81	0.92	0.86	1000
Horse	0.90	0.85	0.88	1000
Ship	0.95	0.85	0.90	1000
Truck	0.92	0.88	0.90	1000
Accuracy			0.84	10,000
Macro Avg	0.84	0.84	0.84	10,000
Weighted Avg	0.84	0.84	0.84	10,000

04

Model Improvement Strategies

Model Improvement Strategies

3 different strategies were attempted to try to improve the CNN model's accuracy.

1

Random Data Augmentations

Random data augmentations can be performed to artificially increase the size and diversity of the training dataset through random flips, rotations, etc.

2

Use RSOA for Hyperparameter Tuning

Hyperparameter tuning with Random Search can be performed to find the most optimal hyperparameters for the training of CNN architecture.

3

Use Reduced Learning Rate on Plateau Scheduler



We can reduce the learning rate by $\frac{1}{4}$ every time the validation loss increases to improve model convergence

Strategies that improved on the original CNN model were implemented into the final implementation of the model.


Model Improvement: Random Data Augmentation

- Random horizontal flips with a 50% probability of occurring were added
- The brightness, contrast, saturation, and hue of image were randomly modified by a small degree
 - Forces the model to learn features that are less reliant on exact color values

Findings:



Findings Cont.:



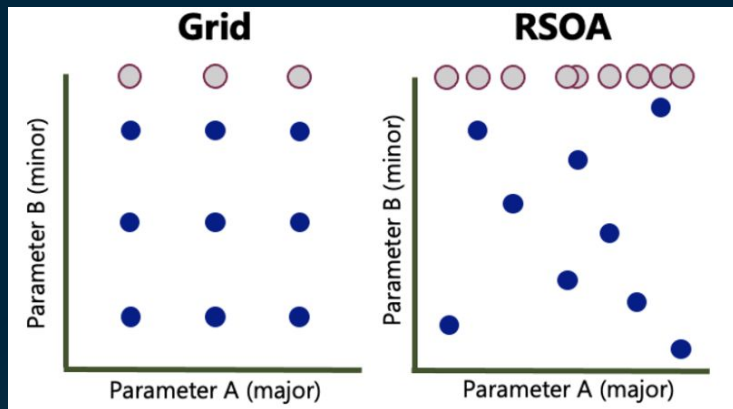
Classification Report				
	Precision	Recall	F1-Score	Support
Airplane	0.85	0.91	0.88	1000
Car	0.96	0.91	0.94	1000
Bird	0.85	0.81	0.82	1000
Cat	0.74	0.90	0.84	1000
Deer	0.79	0.90	0.84	1000
Dog	0.89	0.74	0.81	1000
Frog	0.87	0.90	0.89	1000
Horse	0.91	0.89	0.90	1000
Ship	0.95	0.92	0.93	1000
Truck	0.91	0.95	0.92	1000
Accuracy			0.87	10,000
Macro Avg	0.87	0.87	0.87	10,000
Weighted Avg	0.87	0.87	0.87	10,000

Effects of Random Data Augmentation:

- Less overfitting than before
- Improvement to accuracy by 3% from 84% to **87%**
- Can be implemented into final implementation


Model Improvement: Using RSOA for Hyperparameter Tuning

- The Random Search Optimization Algorithm (RSOA) was used to choose an optimal set of hyperparameters
- 20 random combinations of batch size and learning rate were tested to find optimal hyperparameters
 - Randomly select learning rate between 0.001 and 0.01
 - Randomly select batch size in [32, 64, 128]



RSOA is more effective than other searches since it can cover a wider range of possible parameters (gray circles) in the same number of searches (blue circles)

- Findings:





- A learning rate of 0.0035 and batch size of 32 performed the best for the final accuracy of the CNN architecture on the validation set, but it then only achieved an accuracy of 84% on the unseen testset

Classification Report				
	Precision	Recall	F1 Score	Support
Airplane	0.82	0.91	0.86	1000
Car	0.88	0.95	0.91	1000
Bird	0.80	0.76	0.78	1000
Cat	0.68	0.76	0.72	1000
Deer	0.84	0.81	0.83	1000
Dog	0.85	0.72	0.77	1000
Frog	0.85	0.90	0.87	1000
Horse	0.90	0.86	0.88	1000
Ship	0.95	0.90	0.91	1000
Truck	0.93	0.86	0.89	1000
			0.84	10,000
Accuracy				
Macro Avg	0.84	0.84	0.84	10,000
Weighted Avg	0.84	0.84	0.84	10,000


Model Improvement: Reducing Learning Rate on Plateau Scheduler

- The Reducing LR on Plateau Scheduler was implemented
- LR Reduces by $\frac{1}{4}$ every time that the validation loss increased during the 15 epochs

Findings:



Findings Cont.:



Classification Report				
	Precision	Recall	F1-Score	Support
Airplane	0.85	0.90	0.87	1000
Car	0.96	0.92	0.94	1000
Bird	0.80	0.89	0.81	1000
Cat	0.76	0.74	0.75	1000
Deer	0.85	0.86	0.86	1000
Dog	0.84	0.78	0.81	1000
Frog	0.88	0.91	0.89	1000
Horse	0.91	0.90	0.90	1000
Ship	0.91	0.92	0.92	1000
Truck	0.92	0.92	0.92	1000
Accuracy			0.87	10,000
Macro Avg	0.87	0.87	0.87	10,000
Weighted Avg	0.87	0.87	0.897	10,000

Effects of the Scheduler:

- Improvement to accuracy by 3% from 84% **87%**
- Can be implemented into final implementation


Final Implementation: Combining Random Data Augmentation with Reducing LR on Plateau Scheduler

- Since Strategy 1, Random Data Augmentation and Strategy 3, Reduced Learning Rate on Plateau Scheduler, improved on the original CNN by the greatest margins, a CNN model using both of these strategies was tested

Findings:



Findings Cont.:



Classification Report				
	Precision	Recall	F1-Score	Support
Airplane	0.88	0.92	0.90	1000
Car	0.96	0.95	0.96	1000
Bird	0.82	0.82	0.82	1000
Cat	0.80	0.75	0.77	1000
Deer	0.86	0.89	0.87	1000
Dog	0.85	0.81	0.83	1000
Frog	0.90	0.91	0.90	1000
Horse	0.92	0.92	0.92	1000
Ship	0.92	0.95	0.93	1000
Truck	0.95	0.94	0.94	1000
Accuracy			0.89	10,000
Macro Avg	0.89	0.89	0.89	10,000
Weighted Avg	0.89	0.89	0.89	10,000

- Combining these two strategies yielded in the biggest improvement with an **accuracy of 89%** as compared to only 84% earlier!

Final Conclusion

- The CNN architecture with the Random Data Augmentations and Reduced LR on Plateau Scheduler performed very well with an accuracy of 89% on the CIFAR-10 Dataset!

CNN Model Training Strategy	Accuracy on Test set	Improvement to Original CNN
Original CNN Model	84%	-
CNN Model with Random Data Augmentations	87%	3%
CNN Model with RSOA	84%	0%
CNN Model with Reduced LR on Plateau Scheduler	87%	3%
CNN Model with Random Data Augmentations and Reduced LR on Plateau Scheduler	89%	5%